

Student Name : Bimal Kandel

install or Importing the necessary libraries

1. # Importing the necessary libraries

Import Required Libraries

```
import pandas as pd
import numpy as np
import nltk
from sklearn.model_selection import train_test_split
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, precision_score,
recall_score, f1_score, confusion_matrix
import matplotlib.pyplot as plt
import seaborn as sns
from nltk.corpus import stopwords
from nltk.tokenize import word_tokenize

import nltk
nltk.download('stopwords')      # Stopwords for preprocessing
nltk.download('punkt')        # Tokenizer for text processing
```

```
[nltk_data] Downloading package stopwords to
[nltk_data] /Users/iambimalk/nltk_data...
[nltk_data] Package stopwords is already up-to-date!
[nltk_data] Downloading package punkt to /Users/iambimalk/nltk_data...
[nltk_data] Package punkt is already up-to-date!
```

True

Downloaded the Amazon Customer Reviews dataset from Kaggle and importing them

3. #loading the csv file dataset

```
data = pd.read_csv("//Users/iambimalk/Downloads/7817_1.csv")
```

data

	id	asins	brand	\
0	AVpe7AsMilAPnD_xQ78G	B00QJDU3KY	Amazon	
1	AVpe7AsMilAPnD_xQ78G	B00QJDU3KY	Amazon	
2	AVpe7AsMilAPnD_xQ78G	B00QJDU3KY	Amazon	
3	AVpe7AsMilAPnD_xQ78G	B00QJDU3KY	Amazon	
4	AVpe7AsMilAPnD_xQ78G	B00QJDU3KY	Amazon	
...	
1592	AVpfo9ukilAPnD_xfhuj	B00N08JJZW	Amazon	
1593	AVpfo9ukilAPnD_xfhuj	B00N08JJZW	Amazon	
1594	AVpfo9ukilAPnD_xfhuj	B00N08JJZW	Amazon	
1595	AVpfo9ukilAPnD_xfhuj	B00N08JJZW	Amazon	

1596 AVpfo9ukilAPnD_xfhuj B00N08JJZW Amazon

	categories	colors	\
0	Amazon Devices,mazon.co.uk	NaN	
1	Amazon Devices,mazon.co.uk	NaN	
2	Amazon Devices,mazon.co.uk	NaN	
3	Amazon Devices,mazon.co.uk	NaN	
4	Amazon Devices,mazon.co.uk	NaN	
...	
1592	Amazon Devices & Accessories,Amazon Device Acc...	NaN	
1593	Amazon Devices & Accessories,Amazon Device Acc...	NaN	
1594	Amazon Devices & Accessories,Amazon Device Acc...	NaN	
1595	Amazon Devices & Accessories,Amazon Device Acc...	NaN	
1596	Amazon Devices & Accessories,Amazon Device Acc...	NaN	

	dateAdded	dateUpdated	
dimension \			
0	2016-03-08T20:21:53Z	2017-07-18T23:52:58Z	169 mm x 117 mm x
9.1 mm			
1	2016-03-08T20:21:53Z	2017-07-18T23:52:58Z	169 mm x 117 mm x
9.1 mm			
2	2016-03-08T20:21:53Z	2017-07-18T23:52:58Z	169 mm x 117 mm x
9.1 mm			
3	2016-03-08T20:21:53Z	2017-07-18T23:52:58Z	169 mm x 117 mm x
9.1 mm			
4	2016-03-08T20:21:53Z	2017-07-18T23:52:58Z	169 mm x 117 mm x
9.1 mm			
...	
...			
1592	2016-04-02T14:40:43Z	2017-08-13T08:28:46Z	
NaN			
1593	2016-04-02T14:40:43Z	2017-08-13T08:28:46Z	
NaN			
1594	2016-04-02T14:40:43Z	2017-08-13T08:28:46Z	
NaN			
1595	2016-04-02T14:40:43Z	2017-08-13T08:28:46Z	
NaN			
1596	2016-04-02T14:40:43Z	2017-08-13T08:28:46Z	
NaN			

	ean	keys	...	\
0	NaN	kindlepaperwhite/b00qjdu3ky	...	
1	NaN	kindlepaperwhite/b00qjdu3ky	...	
2	NaN	kindlepaperwhite/b00qjdu3ky	...	
3	NaN	kindlepaperwhite/b00qjdu3ky	...	
4	NaN	kindlepaperwhite/b00qjdu3ky	...	
...	
1592	NaN	alexavoiceremoteforamazonfiretvfiretvstick/b00...	...	
1593	NaN	alexavoiceremoteforamazonfiretvfiretvstick/b00...	...	
1594	NaN	alexavoiceremoteforamazonfiretvfiretvstick/b00...	...	

```

1595 NaN alexavoiceremoteforamazonfiretvfiretvstick/b00... ...
1596 NaN alexavoiceremoteforamazonfiretvfiretvstick/b00... ...

```

	reviews.rating	reviews.sourceURLs
\		
0	5.0	https://www.amazon.com/Kindle-Paperwhite-High-...
1	5.0	https://www.amazon.com/Kindle-Paperwhite-High-...
2	4.0	https://www.amazon.com/Kindle-Paperwhite-High-...
3	5.0	https://www.amazon.com/Kindle-Paperwhite-High-...
4	5.0	https://www.amazon.com/Kindle-Paperwhite-High-...
...
1592	3.0	https://www.amazon.com/Alexa-Voice-Remote-Amaz...
1593	1.0	https://www.amazon.com/Alexa-Voice-Remote-Amaz...
1594	1.0	https://www.amazon.com/Alexa-Voice-Remote-Amaz...
1595	3.0	https://www.amazon.com/Alexa-Voice-Remote-Amaz...
1596	1.0	https://www.amazon.com/Alexa-Voice-Remote-Amaz...

	reviews.text	\
0	I initially had trouble deciding between the p...	
1	Allow me to preface this with a little history...	
2	I am enjoying it so far. Great for reading. Ha...	
3	I bought one of the first Paperwhites and have...	
4	I have to say upfront - I don't like coroporat...	
...
1592	This is not the same remote that I got for my ...	
1593	I have had to change the batteries in this rem...	
1594	Remote did not activate, nor did it connect to...	
1595	It does the job but is super over priced. I fe...	
1596	I ordered this item to replace the one that no...	

	reviews.title
reviews.userCity	\
0	Paperwhite voyage, no regrets!
NaN	
1	One Simply Could Not Ask For More
NaN	
2	Great for those that just want an e-reader
NaN	
3	Love / Hate relationship
NaN	

```

4                                     I LOVE IT
NaN
...
..
1592 I would be disappointed with myself if i produ...
NaN
1593                                     Battery draining remote!!!!
NaN
1594         replacing an even worse remote. Waste of time
NaN
1595                                     Overpriced
NaN
1596 I am sending all of this crap back to amazon a...
NaN

```

	reviews.userProvince	reviews.username	sizes	upc	weight
0	NaN	Cristina M	NaN	NaN	205 grams
1	NaN	Ricky	NaN	NaN	205 grams
2	NaN	Tedd Gardiner	NaN	NaN	205 grams
3	NaN	Dougal	NaN	NaN	205 grams
4	NaN	Miljan David Tanic	NaN	NaN	205 grams
...
1592	NaN	GregAmandawith4	NaN	NaN	4 ounces
1593	NaN	Amazon Customer	NaN	NaN	4 ounces
1594	NaN	Amazon Customer	NaN	NaN	4 ounces
1595	NaN	Meg Ashley	NaN	NaN	4 ounces
1596	NaN	DIANE K	NaN	NaN	4 ounces

```

[1597 rows x 27 columns]

```

Clean the data by removing missing values or duplicates

```

4. #Data Cleaning
data.drop_duplicates(inplace=True)
data.dropna(subset=['reviews.text', 'reviews.rating'], inplace=True)

```

Preprocess the text with tokenization, lowercasing, and removing stop words using NLTK

```

# 5. Preprocessing
def preprocess_text(text):
    tokens = word_tokenize(text.lower()) # Tokenize and lowercase
    stop_words = set(stopwords.words('english'))
    filtered_tokens = [word for word in tokens if word.isalpha() and
word not in stop_words]
    return ' '.join(filtered_tokens)

# Ensure the 'reviews.text' column exists
data = data.rename(columns=lambda x: x.strip()) # Strip column names
to avoid leading/trailing spaces

```

```

if 'reviews.text' in data.columns:
    data['cleaned_review'] =
data['reviews.text'].apply(preprocess_text) # Update 'reviews.text'
to your column name
else:
    raise ValueError("Column 'reviews.text' is missing in the
dataset.")

```

Label sentiment categories as Positive, Negative, or Neutral based on review ratings

```

# 6. Label Sentiment
# Assume rating column has values from 1-5, where 1-2 = Negative, 3 =
Neutral, 4-5 = Positive
def label_sentiment(rating):
    if rating <= 2:
        return 'Negative'
    elif rating == 3:
        return 'Neutral'
    else:
        return 'Positive'

if 'reviews.rating' in data.columns:
    data['sentiment'] = data['reviews.rating'].apply(label_sentiment)
else:
    raise ValueError("Column 'reviews.rating' is missing in the
dataset.")

```

Convert text to numerical features using TF-IDF

```

# 7. Convert Text to Numerical Features
if not data['cleaned_review'].str.strip().replace('',
pd.NA).isna().all():
    vectorizer = TfidfVectorizer(max_features=5000)
    X = vectorizer.fit_transform(data['cleaned_review']).toarray()
    y = data['sentiment']
else:
    raise ValueError("Cleaned reviews result in an empty vocabulary.
Ensure valid data is present.")

```

Split the dataset into training and testing sets with train_test_split

```

# 8. Split Dataset
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

```

Train a classifier using Logistic Regression

9. Train Classifier

```
model = LogisticRegression(max_iter=1000)
model.fit(X_train, y_train)
```

```
LogisticRegression(max_iter=1000)
```

Evaluate the model's performance using metrics like accuracy, precision, recall, and F1-score

10. Evaluate Model

```
y_pred = model.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred, average='weighted')
recall = recall_score(y_test, y_pred, average='weighted')
f1 = f1_score(y_test, y_pred, average='weighted')
conf_matrix = confusion_matrix(y_test, y_pred, labels=['Positive',
'Neutral', 'Negative'])
```

```
/opt/anaconda3/lib/python3.11/site-packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
  _warn_prf(average, modifier, msg_start, len(result))
```

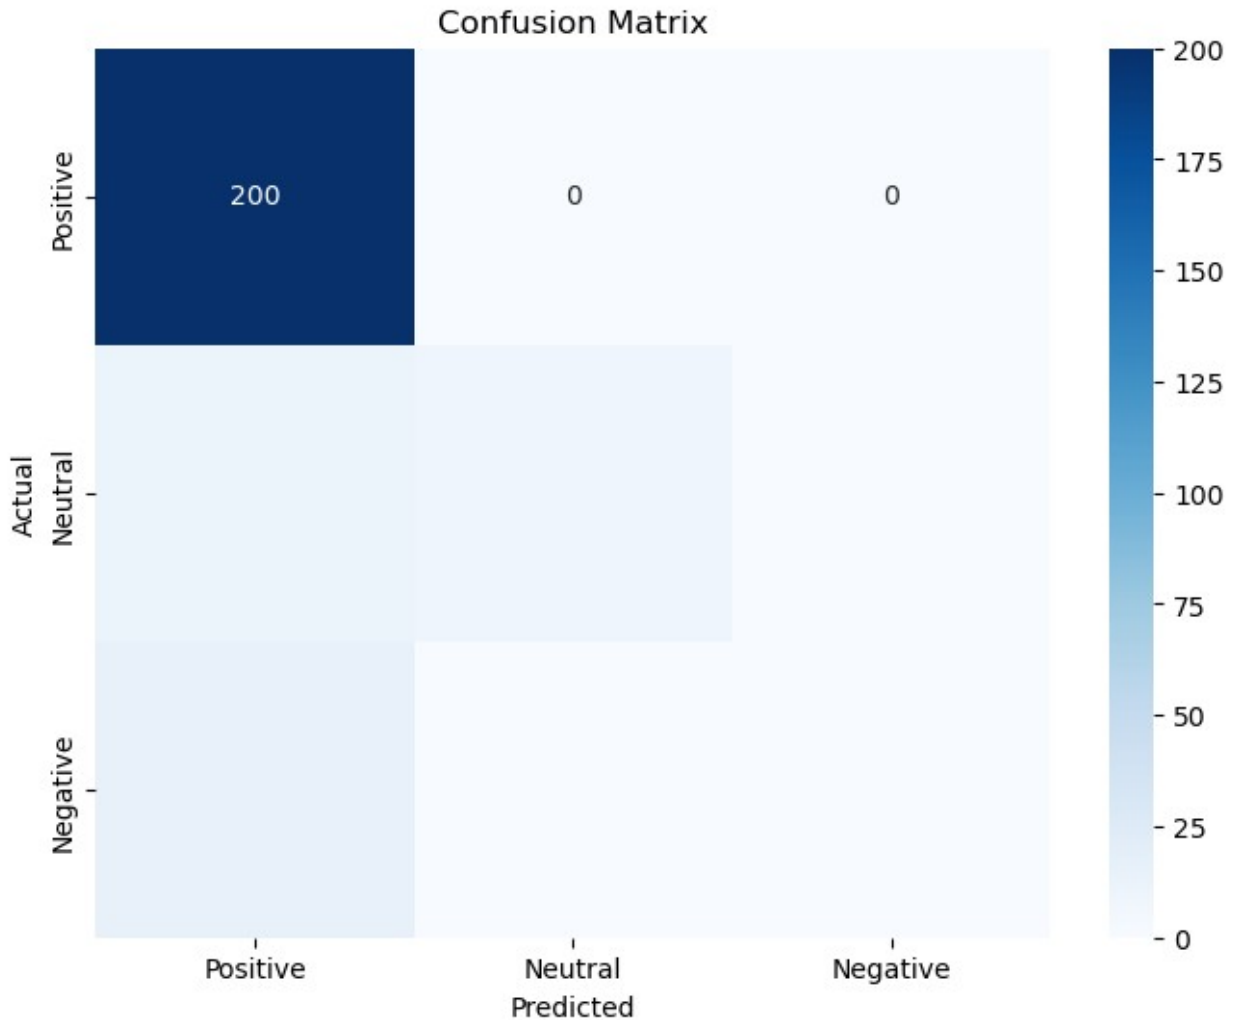
```
print(f"Accuracy: {accuracy}")
print(f"Precision: {precision}")
print(f"Recall: {recall}")
print(f"F1-Score: {f1}")
```

```
Accuracy: 0.8898305084745762
Precision: 0.8389455527223638
Recall: 0.8898305084745762
F1-Score: 0.8531432810458524
```

Visualize results with confusion matrices and sentiment distribution plots

11. Visualize Results

```
plt.figure(figsize=(8, 6))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues',
xticklabels=['Positive', 'Neutral', 'Negative'],
yticklabels=['Positive', 'Neutral', 'Negative'])
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix')
plt.show()
```



```
# Sentiment distribution
if 'Sentiment' in data.columns:
    sentiment_counts = data['Sentiment'].value_counts()
    plt.figure(figsize=(6, 4))
    sentiment_counts.plot(kind='bar', color=['green', 'blue', 'red'])
    plt.title('Sentiment Distribution')
    plt.xlabel('Sentiment')
    plt.ylabel('Count')
    plt.show()
else:
    print("Column 'Sentiment' does not exist.")
```

Column 'Sentiment' does not exist.

```
14.# Save Results
data[['reviews.text', 'cleaned_review',
'sentiment']].to_csv('processed_reviews.csv', index=False)
```